Project Report
(IO Tracing)

Asif Ali
143059009
CSE Mtech 1

Problem Statement

- To know about the operations of each I/O layer performed when a process request for I/O from a block device.

Architecture

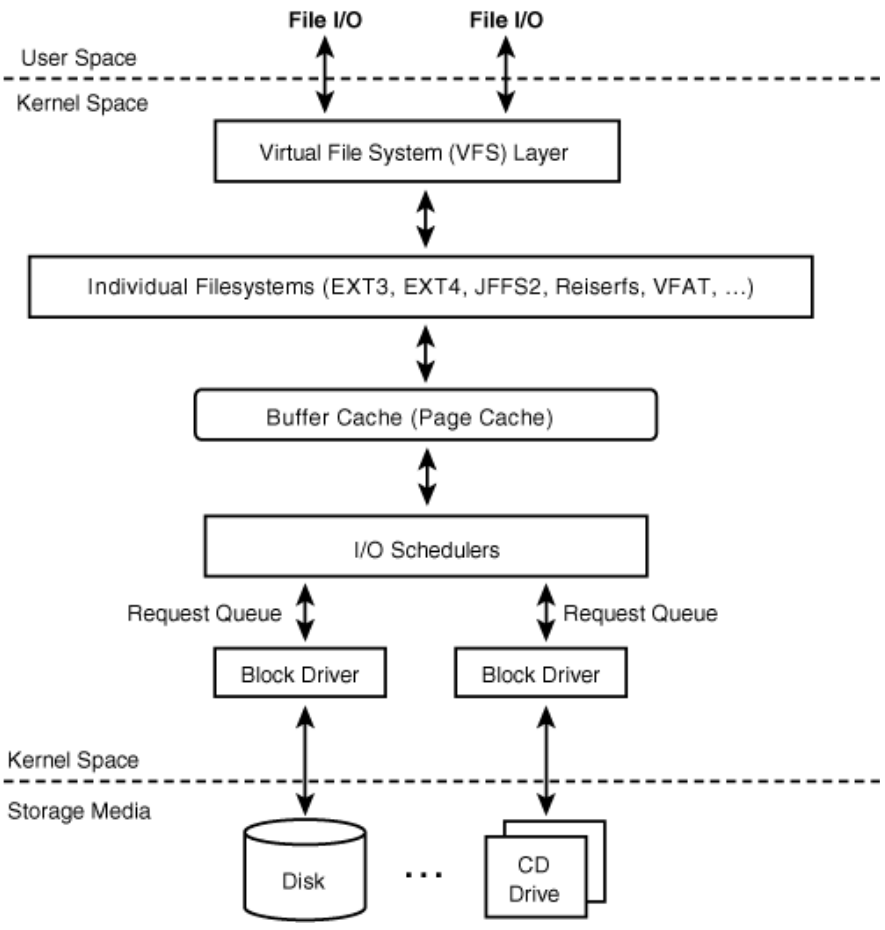

An Example: Blktrace

Introduction

Blktrace is a block layer IO tracing mechanism which provides detailed information about request queue operations up to user space. There are three major components: a kernel component, a utility to record the i/o trace information for the kernel to user space, and utilities to analyse and view the trace information. This man page describes blktrace, which records the i/o event trace information for a specific block device to a file.
The blktrace utility extracts event traces from the kernel (via the relaying through the debug file system).

blktrace – live

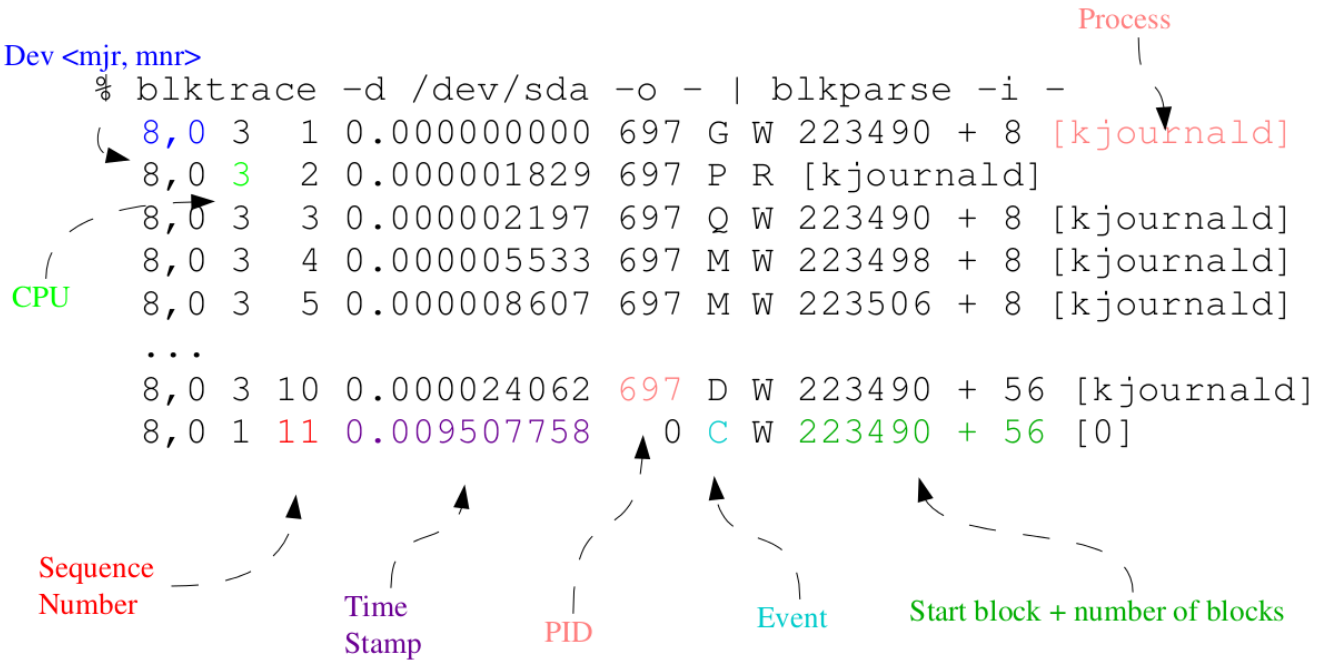

```
CPU1 (8,0):
 Reads Queued:           0,        0KiB  Writes Queued:           7,      128KiB
 Read Dispatches:        0,        0KiB  Write Dispatches:        7,      128KiB
 Reads Completed:        0,        0KiB  Writes Completed:       11,      168KiB
 Read Merges:            0              Write Merges:           25
 IO unplugs:             0              Timer unplugs:           0
...
CPU3 (8,0):
 Reads Queued:           0,        0KiB  Writes Queued:           1,       28KiB
 Read Dispatches:        0,        0KiB  Write Dispatches:        1,       28KiB
 Reads Completed:        0,        0KiB  Writes Completed:        0,        0KiB
 Read Merges:            0              Write Merges:            6
 IO unplugs:             0              Timer unplugs:           0

Total (8,0):
 Reads Queued:           0,        0KiB  Writes Queued:          11,      168KiB
 Read Dispatches:        0,        0KiB  Write Dispatches:       11,      168KiB
 Reads Completed:        0,        0KiB  Writes Completed:       11,      168KiB
 Read Merges:            0              Write Merges:           31
 IO unplugs:             0              Timer unplugs:           3

Events (8,0): 89 entries, 0 skips
```

Architecture of Blktrace

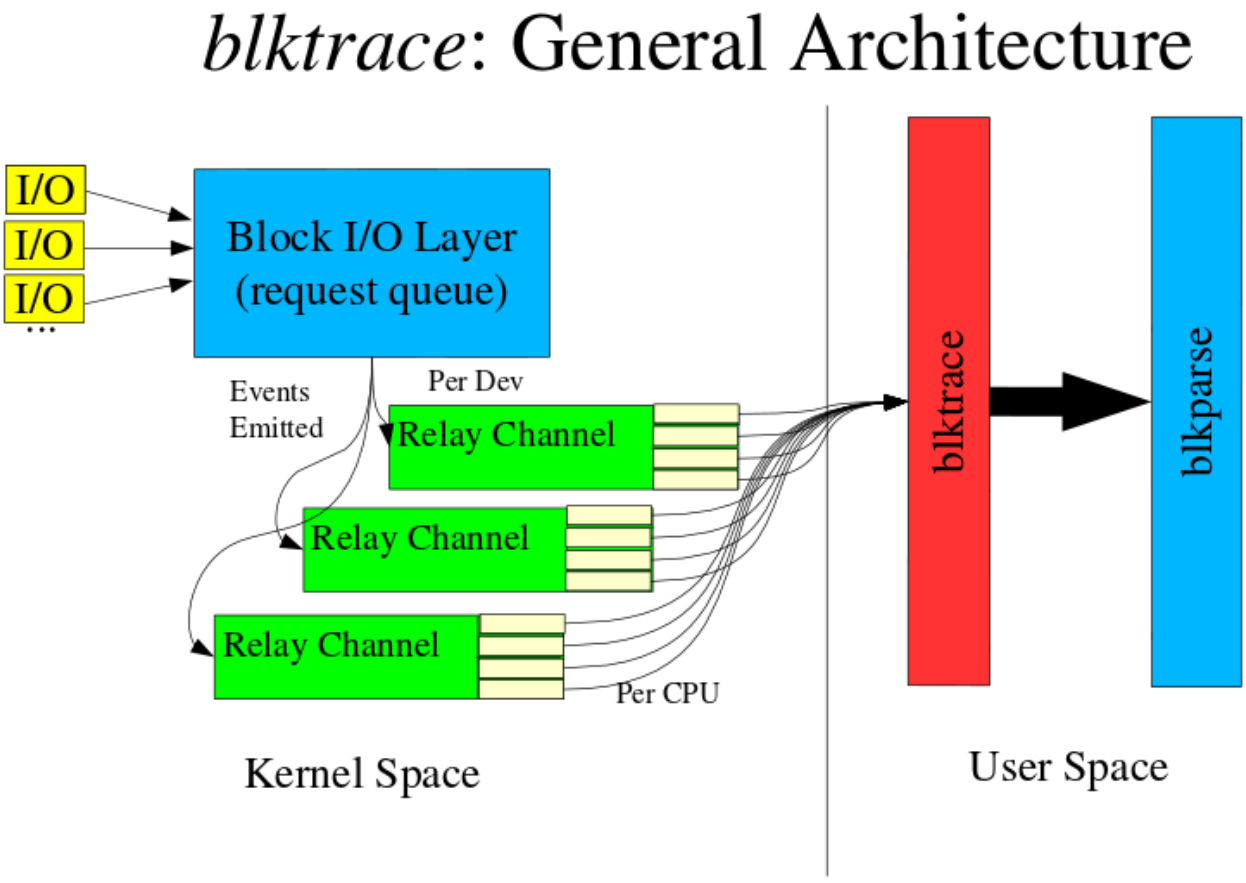

Components of My Project:

- Tracing Module
- Parse script (written in python)

Methodology And Approach

Put three hook functions(block_fun, block_comp, block_requeue) in generic_make_request function, blk_requeue_request, blk_put_request function respectively. Then pass the state of request_queue to our module.Our module just prints that states (traces).
What our parse script does is to take that traces from /var/log/syslog file and collect these trace and make combined data.

Experiments:

What we did is to run our trace program and see the result.

```
[ 5953.405491] Major 8 minor 1
[ 5953.405500] disk name sda
[ 5954.528792] 8,1  1  -5950.-874938444 'C'
[ 5954.529558] 8,1  0  2  1.1240047142  919  I  R  0+8  jbd2/sda1-8
[ 5954.529626] 8,1  0  3  1.1241189570  919  I  R  0+8  jbd2/sda1-8
[ 5954.529634] 8,1  0  4  1.1241285700  919  I  R  0+8  jbd2/sda1-8
[ 5954.529644] 8,1  0  5  1.1241338630  919  I  R  0+8  jbd2/sda1-8
[ 5954.529648] 8,1  0  6  1.1241386300  919  I  R  0+8  jbd2/sda1-8
[ 5954.529653] 8,1  0  7  1.1241433540  919  I  R  0+8  jbd2/sda1-8
[ 5954.529658] 8,1  0  8  1.1241481630  919  I  R  0+8  jbd2/sda1-8
[ 5954.529663] 8,1  0  9  1.1241528670  919  I  R  0+8  jbd2/sda1-8
[ 5954.529668] 8,1  0  10  1.1241576230  919  I  R  0+8  jbd2/sda1-8
[ 5954.529673] 8,1  0  11  1.1241628770  919  I  R  0+8  jbd2/sda1-8
[ 5954.529678] 8,1  0  12  1.1241678690  919  I  R  0+8  jbd2/sda1-8
[ 5954.529682] 8,1  0  13  1.1241726130  919  I  R  0+8  jbd2/sda1-8
[ 5954.529687] 8,1  0  14  1.1241773540  919  I  R  0+8  jbd2/sda1-8
[ 5954.529692] 8,1  0  15  1.1241821060  919  I  R  0+8  jbd2/sda1-8
[ 5954.529697] 8,1  0  16  1.1241868340  919  I  R  0+8  jbd2/sda1-8
[ 5954.529702] 8,1  0  17  1.1241915580  919  I  R  0+8  jbd2/sda1-8
[ 5954.529708] 8,1  0  18  1.1241965300  919  I  R  0+8  jbd2/sda1-8
[ 5954.529713] 8,1  0  19  1.1242027960  919  I  R  0+8  jbd2/sda1-8
[ 5954.529718] 8,1  0  20  1.1242078660  919  I  R  0+8  jbd2/sda1-8
[ 5954.529725] 8,1  0  21  1.1242128010  919  I  R  0+8  jbd2/sda1-8
[ 5954.529730] 8,1  0  22  1.1242198800  919  I  R  0+8  jbd2/sda1-8
[ 5954.530348] 8,1  0  23  1.1242249430  919  I  R  0+8  jbd2/sda1-8
[ 5956.576777] 8,1  25  1.1248407050  'C'
[ 5958.624769] 8,1  26  1.1248407050  'C'
[ 5959.712202] 8,1  0  27  7.-693321940  919  I  R  0+8  jbd2/sda1-8
[ 5959.712232] 8,1  0  28  7.-693280665  919  I  R  0+8  jbd2/sda1-8
[ 5959.712714] 8,1  0  29  7.-692799639  919  I  R  0+8  jbd2/sda1-8
[ 5960.672837] 8,1  30  7.-692799639  'C'
[ 5962.720840] 8,1  31  7.-692799639  'C'
[ 5964.768834] 8,1  32  7.-692799639  'C'
[ 5964.769634] 8,1  0  33  12.-635889269  919  I  R  0+8  jbd2/sda1-8
```

Firefox Web Browser

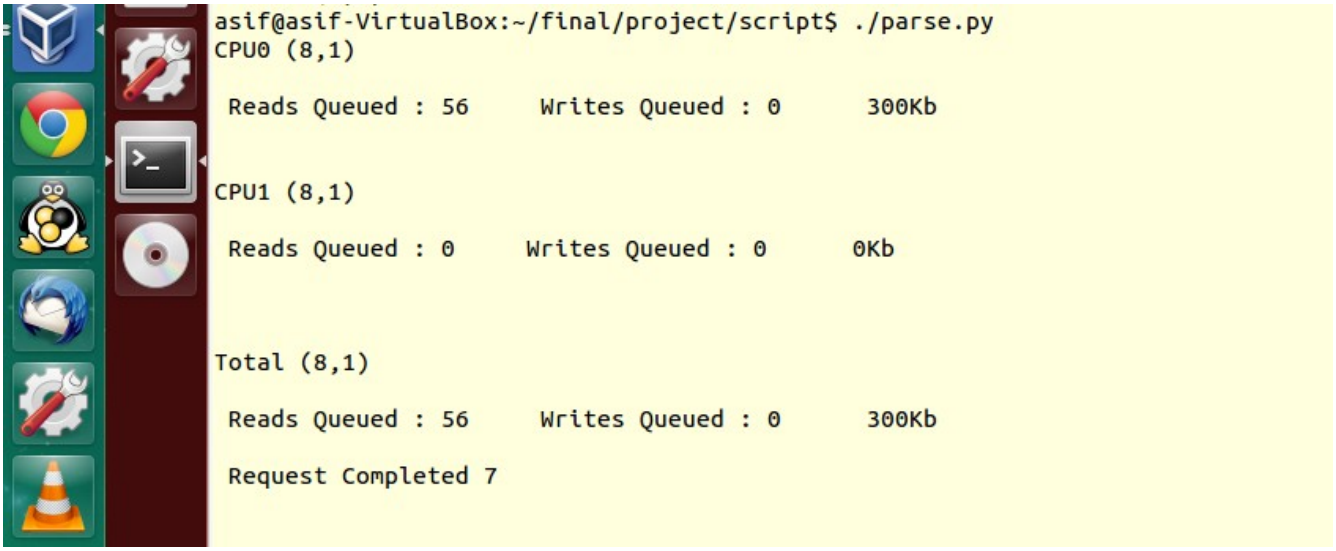

We did'nt perform any write request that's there si no writing traces came in our experiment.

References:

1.http://www.embeddedlinux.org.cn/EssentialLinuxDeviceDrivers/final/ch14lev1sec2.html

2.Blktrace User Guide

http://www.cse.unsw.edu.au/~aaronc/iosched/doc/blktrace.html
3.Block I/O Tracing
http://www.mimuw.edu.pl/~lichota/09-10/Optymalizacja-open-source/Materialy/10%20-%20Dysk/gelato_ICE06apr_blktrace_brunelle_hp.pdf
4.Makelinux.net
5.Reading block I/O from kernel module
http://stackoverflow.com/questions/14628275/reading-block-device-from-kernel-3-7-module-segfault-in-submit-bio-bd-disk-i/15946581#15946581